

EdgeCount: Statistical Theory and Algorithmic Efficiency

Joel R. Pradines

December 6, 2025

Abstract

This vignette details the statistical framework underlying the **EdgeCount** package. First, the Random Graph with Given Expected Degrees (RGGED) model is described. Second, the derivation of edge-count statistics is presented. Third, the algebraic optimizations that allow the package to calculate connectedness statistics with linear time complexity $O(N)$ are explained. Last, the implementation of vectorized functions contained in the package is discussed, and it is shown that they have optimal time complexity.

1 Introduction

The **EdgeCount** package provides methods to analyze the connectedness of vertex sets in large, sparse, undirected graphs. Whether analyzing protein-protein interaction networks or bipartite term-element graphs, the fundamental question is: *Is the observed number of edges, called an edge count, within a set (or between two sets) significantly higher than expected by chance?*

To answer this, the concept of *chance* relies on a null model of random graph that preserve the expected degrees of vertices (Chung and Lu, 2002a,b); that is, the average degree of a vertex over all random graph realizations is its observed value in the original network. Taking into account vertex degrees is important, because they often vary by orders of magnitude. The analytical expressions for the likelihoods of edge-count variables in this model were derived by Pradines et al. (2005).

2 The Random Graph Model

Let $G(V, E)$ be an undirected graph with no loops and no multiple edges.

- Let $N = |V|$ be the order of the graph (number of vertices).
- Let $M = |E|$ be the size of the graph (number of edges).
- Let k_i be the degree of vertex v_i .

To determine statistical significance, we compare observed edge counts in G to a null model: the ensemble of random graphs that preserve the expected degree sequence of G . In the Random Graph with Given Expected Degrees (RGGED) model, edges are treated as independent Bernoulli random variables.

The probability p_{ij} of observing an edge between vertices v_i and v_j is given by Pradines et al. (2005):

$$p_{ij} = \min \left(1, \frac{k_i k_j}{2M} \right) \quad (1)$$

2.1 Fast Approximation for Sparse Graphs

For sparse graphs where $\max(k_i k_j) \ll 2M$, this approximates the probability of edge formation in random graphs with fixed degrees. When the graph is sufficiently sparse (as determined by the package function `summarize_suitability_fast`), we can use the approximation:

$$p_{ij} \approx \frac{k_i k_j}{2M} \quad (2)$$

This approximation is central to the algorithmic efficiency of the package, as it allows for the factorization of degree sums.

3 Edge Count Probabilities

Let Z be a random variable representing the count of edges in a specific subgraph configuration (e.g., edges within a set). Z is the sum of independent Bernoulli trials. While the exact distribution of Z is a Poisson-Binomial distribution, it is well-approximated by a Poisson distribution (or a truncated Poisson distribution) when probabilities p_{ij} are small (Pradines et al., 2005).

The probability of observing z edges is approximated by:

$$P(Z = z) = \frac{e^{-\lambda} \lambda^z}{\alpha z!} \quad (3)$$

where:

- $\lambda = \mathbb{E}[Z] = \sum p_{ij}$ is the expected number of edges.
- $\alpha = e^{-\lambda} \sum_{h=0}^m \frac{\lambda^h}{h!}$ is the normalization factor accounting for the maximum possible number of edges m in the subgraph.

The p -value, or the likelihood of observing x or more edges, is:

$$p\text{-value} = P(Z \geq x) = \frac{e^{-\lambda}}{\alpha} \sum_{h=x}^m \frac{\lambda^h}{h!} \quad (4)$$

4 Algorithmic Optimization

A key feature of **EdgeCount** is its speed. Calculating λ involves summing probabilities over pairs of vertices. A naive implementation requires iterating over all pairs, resulting in quadratic time complexity. However, by utilizing the "Fast" approximation (Equation 2), we can factorize these sums to achieve linear time complexity.

4.1 Within-Set Connectedness

Consider a set of vertices $S \subseteq V$ with size $n = |S|$. We wish to calculate the expected number of edges λ_{in} existing internally among members of S .

$$\lambda_{in} = \sum_{\{i,j\} \subseteq S, i < j} p_{ij} \approx \sum_{i < j} \frac{k_i k_j}{2M} \quad (5)$$

A naive summation requires $O(n^2)$ operations. To optimize this, we recall the algebraic identity for the square of a sum:

$$\left(\sum_{i \in S} k_i \right)^2 = \sum_{i \in S} k_i^2 + 2 \sum_{i < j} k_i k_j \quad (6)$$

Rearranging for the cross-terms:

$$\sum_{i < j} k_i k_j = \frac{1}{2} \left[\left(\sum_{i \in S} k_i \right)^2 - \sum_{i \in S} k_i^2 \right] \quad (7)$$

Substituting this into the expression for λ_{in} :

$$\lambda_{in} \approx \frac{1}{2M} \cdot \frac{1}{2} \left[\left(\sum_{i \in S} k_i \right)^2 - \sum_{i \in S} k_i^2 \right] = \frac{1}{4M} \left[\left(\sum_{i \in S} k_i \right)^2 - \sum_{i \in S} k_i^2 \right] \quad (8)$$

Complexity: We now only need to compute $\sum k_i$ and $\sum k_i^2$. This requires iterating through S once, reducing the complexity to $O(n)$. This optimization is implemented in the package function `calculate_in_stats_fast_vectorized`.

4.2 Between-Set Connectedness

Consider two disjoint sets of vertices S_1 and S_2 . We wish to calculate the expected number of edges λ_{bt} bridging these two sets.

$$\lambda_{bt} = \sum_{i \in S_1} \sum_{j \in S_2} p_{ij} \approx \sum_{i \in S_1} \sum_{j \in S_2} \frac{k_i k_j}{2M} \quad (9)$$

A naive summation requires $O(|S_1| \cdot |S_2|)$ operations. Because the terms k_i and k_j are independent in the product, we can factor the sum:

$$\lambda_{bt} \approx \frac{1}{2M} \left(\sum_{i \in S_1} k_i \right) \left(\sum_{j \in S_2} k_j \right) \quad (10)$$

Complexity: We sum the degrees in S_1 ($O(|S_1|)$) and S_2 ($O(|S_2|)$) separately and multiply the results. The total complexity is $O(|S_1| + |S_2|)$, which is linear relative to the input size. This optimization is implemented in `calculate_between_stats_fast_vectorized`.

5 Effect Size: The Anscombe Ratio

While the p -value indicates significance, it is dependent on sample size. To measure the magnitude of connectedness, **EdgeCount** employs the log Anscombe Ratio (lAr), which stabilizes the variance of Poisson-distributed variables.

$$lAr = \frac{1}{2} \log_2 \left(\frac{x + 3/8}{\lambda + 3/8} \right) \quad (11)$$

where x is the observed edge count and λ is the expected edge count.

- $lAr > 0$: The set is more connected than expected (enrichment).
- $lAr < 0$: The set is less connected than expected (depletion).

The addition of $3/8$ is an optimal correction for bias in the log transformation of Poisson variables (Pradines et al., 2020).

6 Vertex Set Enrichment Analysis (VSEA)

The VSEA method implemented in `run_vsea` adapts the principles of Gene Set Enrichment Analysis (Subramanian et al., 2005) to graph topology.

Standard GSEA uses a Kolmogorov-Smirnov-like running sum statistic. In VSEA, the "local" correlation at rank i is replaced by the edge-count probability framework derived by Pradines et al. (2005). As we traverse the ranked list of elements, at step i (considering the top i elements as set S_i), we calculate:

1. The observed edges x_i between S_i and a target Term T .
2. The expected edges λ_i using the RGGED model.
3. The running score based on the Log Anscombe Ratio.

Crucially, because λ_i accounts for the degrees of the elements in S_i , **VSEA inherently penalizes hubs**. If the top-ranked elements are merely "sticky" (high degree) but not specifically connected to Term T , λ_i will be very high, reducing the lAr score. This distinguishes VSEA from standard hypergeometric enrichment which only counts overlaps, ignoring the varying likelihood of connections for high-degree elements.

References

- Fan Chung and Linyuan Lu. Connected components in random graphs with given expected degree sequences. Annals of Combinatorics, 6(2):125–145, 2002a. doi: 10.1007/PL00012580.*
- Fan R. K. Chung and Linyuan Lu. The average distances in random graphs with given expected degrees. Proceedings of the National Academy of Sciences of the United States of America, 99:15879–15882, 2002b. URL <https://api.semanticscholar.org/CorpusID:17723054>.*
- J Pradines, V Farutin, S Rowley, and V Dancik. Analyzing protein lists with large networks: edge-count probabilities in random graphs with given expected degrees. Journal of Computational Biology, 12(2):113–128, 2005. URL <https://joelrmath.github.io/pdfs/JCB.pdf>.*
- JR Pradines, N Cilfone, A Ghavami, V Farutin, E Kurtagic, J Guess, A Manning, and I Capila. Enhancing reproducibility of gene-expression analysis with known protein functional relationships: the concept of well-associated protein. PLOS Computational Biology, 16(2):e1007684, 2020. URL <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1007684>.*
- A Subramanian, P Tamayo, VK Mootha, S Mukherjee, BL Ebert, MA Gillette, A Paulovich, SL Pomeroy, TR Golub, ES Lander, and JP Mesirov. Gene set enrichment analysis: a knowledge-based approach for interpreting genome-wide expression profiles. Proceedings of the National Academy of Sciences, 102(43):15545–15550, 2005. doi: 10.1073/pnas.0506580102.*